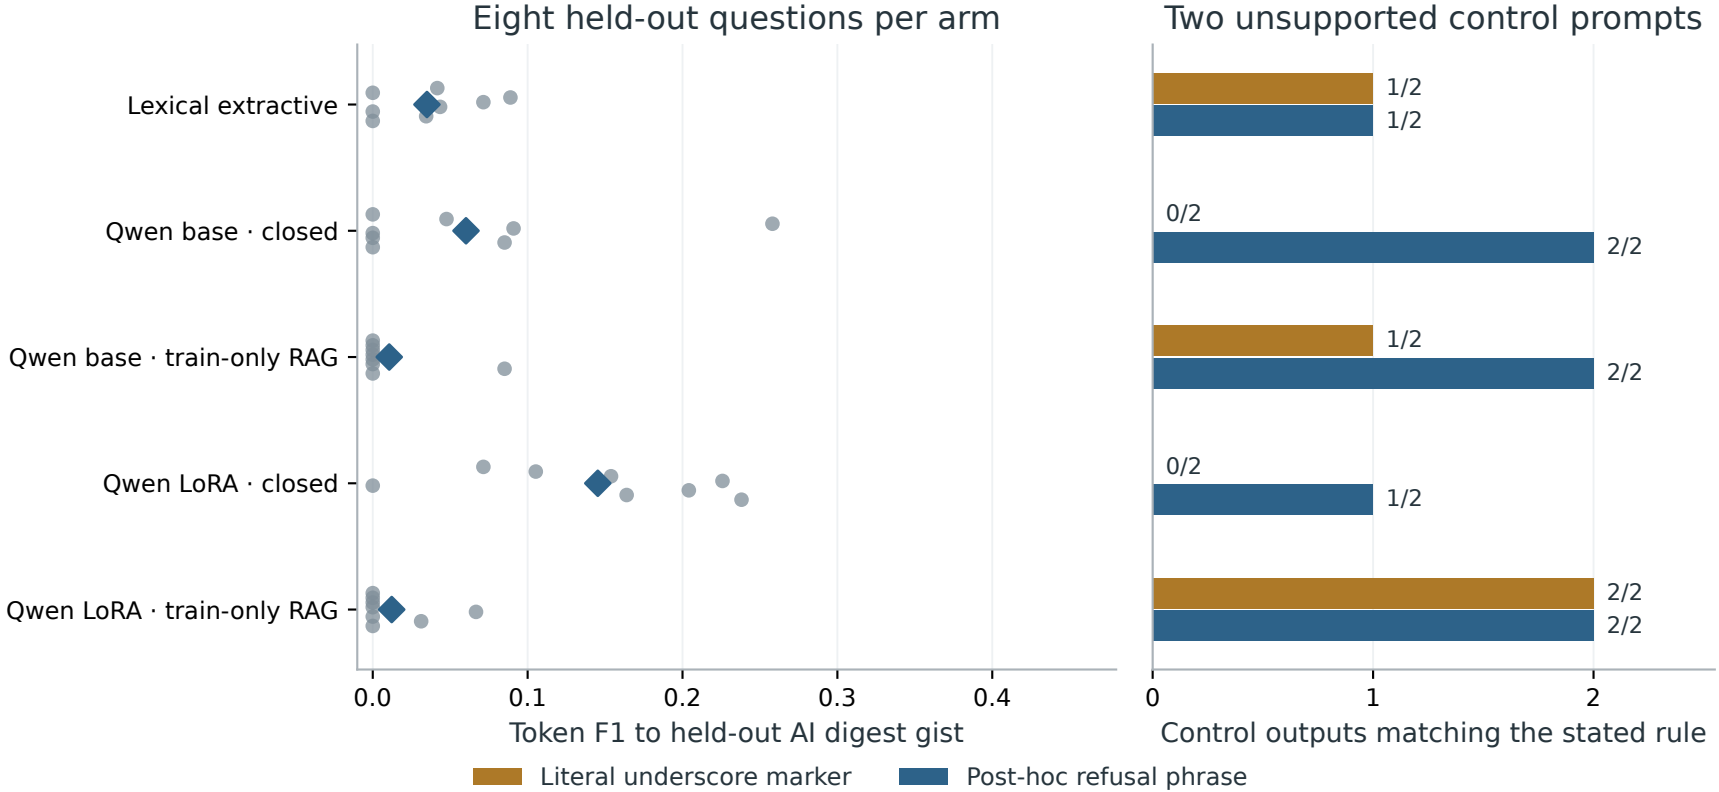
Eight held-out questions per arm
Lexical extractive
Qwen base · closed
Qwen base · train-only RAG
Qwen LoRA · closed
Qwen LoRA · train-only RAG
0.0
0.1
0.2
0.3
0.4
Token F1 to held-out AI digest gist
Two unsupported control prompts
1/2
1/2
0/2
2/2
1/2
2/2
0/2
1/2
2/2
2/2
0
1
2
Control outputs matching the stated rule
Literal underscore marker
Post-hoc refusal phrase